

# Homework 7-8

November 14, 2023

## 1 Homework 7-8

### 1.1 Regression

Load the data.

```
[ ]: import pandas as pd

df = pd.read_csv('Advertising.csv')
df = df.drop('Unnamed: 0', axis=1)
df.head(5)
```

```
[ ]:      TV  radio  newspaper  sales
0  230.1   37.8      69.2    22.1
1   44.5   39.3      45.1    10.4
2   17.2   45.9      69.3     9.3
3  151.5   41.3      58.5    18.5
4  180.8   10.8      58.4    12.9
```

Scale the features.

```
[ ]: from sklearn.preprocessing import MinMaxScaler
import numpy as np

x = np.array(df[['TV', 'radio', 'newspaper']])
y = np.array(df[['sales']])

x_scaled = MinMaxScaler().fit_transform(x)
```

Split scaled data into training and testing splits.

```
[ ]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(
    x_scaled, y, test_size=.2, random_state=42
)
```

#### 1.1.1 Exercise 1

Fit linear regression to training data.

```
[ ]: from sklearn.linear_model import LinearRegression

ex1_model = LinearRegression().fit(x_train, y_train)
```

Read off the learned model parameters.

```
[ ]: print(f'Coefficients: {ex1_model.coef_}')
      print(f'Intercept: {ex1_model.intercept_}')
```

```
Coefficients: [[13.22651832  9.38407469  0.3139387 ]]
Intercept: [3.01120633]
```

Compare  $R^2$  scores of training and test data.

```
[ ]: print(f'Training R^2: {ex1_model.score(x_train, y_train)}')
      print(f'Test R^2: {ex1_model.score(x_test, y_test)}')
```

```
Training R^2: 0.8957008271017817
Test R^2: 0.8994380241009119
```

The training and testing  $R^2$  are both fairly close to 1 and similar in value (though for different values of `random_state`, that is, different training and testing splits, the test  $R^2$  tends to be a bit lower than the training  $R^2$ ).

### 1.1.2 Exercise 2

Fit  $K$ -nearest neighbors model to the training data with  $K = 1$ .

```
[ ]: from sklearn.neighbors import KNeighborsRegressor

ex2_model = KNeighborsRegressor(n_neighbors=1).fit(x_train, y_train)
```

Compare  $R^2$  scores for training and testing data.

```
[ ]: print(f'Training R^2: {ex2_model.score(x_train, y_train)}')
      print(f'Test R^2: {ex2_model.score(x_test, y_test)}')
```

```
Training R^2: 1.0
Test R^2: 0.9025380308603167
```

$K$ -nearest neighbors is perfect ( $R^2 = 1$ ) on the training data when  $K = 1$  by construction. On the test data, however, its performance is similar to that of the linear regression in Exercise 1.

Repeat the above for  $K \in \{2, 3, \dots, 10\}$ .

```
[ ]: ex2_models = [ex2_model]
      for k in range(2, 11):
          ex2_models.append(KNeighborsRegressor(k).fit(x_train, y_train))
```

Print  $R^2$  scores from each model.

```
[ ]: for k in range(10):
      print(f'K = {k + 1}')
      print(f'Training R^2: {ex2_models[k].score(x_train, y_train)}')
      print(f'Test R^2: {ex2_models[k].score(x_test, y_test)}')
      print()
```

```
K = 1
Training R^2: 1.0
Test R^2: 0.9025380308603167

K = 2
Training R^2: 0.9812430719552749
Test R^2: 0.9324894681867459

K = 3
Training R^2: 0.9718620977717801
Test R^2: 0.9364836092038614

K = 4
Training R^2: 0.9671987204202612
Test R^2: 0.9299073794472469

K = 5
Training R^2: 0.9665527977251915
Test R^2: 0.9337192077502264

K = 6
Training R^2: 0.9632464215142899
Test R^2: 0.9498780845328242

K = 7
Training R^2: 0.9597421955914276
Test R^2: 0.9503836147120786

K = 8
Training R^2: 0.9570375379536363
Test R^2: 0.9461417178612622

K = 9
Training R^2: 0.9512721124742588
Test R^2: 0.9420403901172719

K = 10
Training R^2: 0.9477586389705528
Test R^2: 0.9382764359650904
```

For  $K > 1$ , the  $K$ -nearest neighbors regression is no longer perfect on the training data. Indeed, as

$K$  increases the training  $R^2$  decreases, but the test  $R^2$  tends to increase (though it is still always lower than the training  $R^2$ ).

### 1.1.3 Exercise 3

Create linear model using keras.

```
[ ]: import tensorflow as tf

ex3_model = tf.keras.Sequential(layers=[
    tf.keras.layers.Input(3), tf.keras.layers.Dense(1)
], name='Exercise3Model')

optim = tf.keras.optimizers.Adam(learning_rate=.15)
ex3_model.compile(loss='mse', optimizer=optim)
```

Train with 5000 epochs.

```
[ ]: x_train_t = tf.convert_to_tensor(x_train)
y_train_t = tf.convert_to_tensor(y_train)

ex3_model.fit(
    x_train_t, y_train_t,
    epochs=5000,
    batch_size=len(x_train),
    verbose=0
)
```

```
[ ]: <keras.src.callbacks.History at 0x2b15fef9610>
```

Get the trained parameters.

```
[ ]: print(f'Weights: {ex3_model.layers[0].weights[0].value()}')
print(f'Bias: {ex3_model.layers[0].bias.value()}')
```

```
Weights: [[13.226515 ]
 [ 9.384072 ]
 [ 0.31393996]]
Bias: [3.0112088]
```

These parameters are comparable to those obtained using ordinary least squares (Exercise 1).

```
[ ]: from sklearn.metrics import r2_score

print(f'Training R^2: {r2_score(y_train, ex3_model.predict(x_train))}')
print(f'Test R^2: {r2_score(y_test, ex3_model.predict(x_test))}')
```

```
5/5 [=====] - 0s 499us/step
5/5 [=====] - 0s 499us/step
Training R^2: 0.8957008421597463
```

```
2/2 [=====] - 0s 1ms/step
Test R^2: 0.8994379998671596
```

The  $R^2$  values are similar for the training and test splits and similar to the  $R^2$  values obtained in Exercise 1.

Overall, this model is nearly the same as that obtained in Exercise 1.

#### 1.1.4 Exercise 4

Create the model.

```
[ ]: ex4_model = tf.keras.Sequential([
    tf.keras.layers.Input(3),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(1)
], name='Exercise4Model')

optim = tf.keras.optimizers.Adam(learning_rate=.05)
ex4_model.compile(loss='mse', optimizer=optim)
```

Train for 5000 epochs.

```
[ ]: ex4_model.fit(
    x_train_t, y_train_t,
    batch_size=len(x_train), verbose=0,
    epochs=5000
)
```

```
[ ]: <keras.src.callbacks.History at 0x2b161490590>
```

Get  $R^2$  scores for training and testing data.

```
[ ]: print(f'Training R^2: {r2_score(y_train_t, ex4_model.predict(x_train_t,
    ↪ verbose=0))}')
print(f'Test R^2: {r2_score(y_test, ex4_model.predict(x_test, verbose=0))}')
```

```
Training R^2: 0.9969724869682755
```

```
Test R^2: 0.988450162299337
```

This model performs much better on both training and testing data than any of the previous ones.

## 1.2 Classification

Load the data, scale features, and split into training and testing data.

```
[ ]: df = pd.read_csv('diagnosis.csv')
df.head(5)
```

```
[ ]:      id diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  \
0      842302         M        17.99         10.38         122.80        1001.0
1      842517         M        20.57         17.77         132.90        1326.0
2      84300903        M        19.69         21.25         130.00        1203.0
3      84348301         M        11.42         20.38          77.58         386.1
4      84358402         M        20.29         14.34         135.10        1297.0

      smoothness_mean  compactness_mean  concavity_mean  concave points_mean  \
0          0.11840         0.27760         0.3001         0.14710
1          0.08474         0.07864         0.0869         0.07017
2          0.10960         0.15990         0.1974         0.12790
3          0.14250         0.28390         0.2414         0.10520
4          0.10030         0.13280         0.1980         0.10430

      ... texture_worst  perimeter_worst  area_worst  smoothness_worst  \
0      ...          17.33          184.60        2019.0          0.1622
1      ...          23.41          158.80        1956.0          0.1238
2      ...          25.53          152.50        1709.0          0.1444
3      ...          26.50           98.87         567.7          0.2098
4      ...          16.67          152.20        1575.0          0.1374

      compactness_worst  concavity_worst  concave points_worst  symmetry_worst  \
0          0.6656         0.7119         0.2654         0.4601
1          0.1866         0.2416         0.1860         0.2750
2          0.4245         0.4504         0.2430         0.3613
3          0.8663         0.6869         0.2575         0.6638
4          0.2050         0.4000         0.1625         0.2364

      fractal_dimension_worst  Unnamed: 32
0          0.11890          NaN
1          0.08902          NaN
2          0.08758          NaN
3          0.17300          NaN
4          0.07678          NaN
```

[5 rows x 33 columns]

Get relevant features.

```
[ ]: # output class is 1 if diagnosis is malignant, 0, if benign
y = np.where(df['diagnosis'] == 'M', 1., 0.)
x = np.array(df[['radius_mean', 'texture_mean', 'smoothness_mean']])
```

Scale features using MinMaxScaler.

```
[ ]: x_scaled = MinMaxScaler().fit_transform(x)
```

Split data into train and test splits.

```
[ ]: x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=.2,
↳ random_state=42)
```

### 1.2.1 Exercise 5

Perform logistic regression on the training data.

```
[ ]: from sklearn.linear_model import LogisticRegression

ex5_model = LogisticRegression(penalty=None).fit(x_train, y_train)
```

Print learned parameters.

```
[ ]: print(f'Coefficients: {ex5_model.coef_}')
print(f'Intercept: {ex5_model.intercept_}')
```

Coefficients: [[27.33747378 10.76746566 15.43661982]]  
Intercept: [-19.74212166]

Compute accuracy on training and testing data.

```
[ ]: print(f'Training accuracy: {ex5_model.score(x_train, y_train)*100 : .02f}%')
print(f'Test accuracy: {ex5_model.score(x_test, y_test)*100 : .02f}%')
```

Training accuracy: 92.75%

Test accuracy: 94.74%

The training and test accuracies are both pretty good (>90%) and are similar. By chance, we got slightly better testing accuracy than training accuracy.

### 1.2.2 Exercise 6

Create and train a  $K$ -nearest neighbors model with  $K = 1$ .

```
[ ]: from sklearn.neighbors import KNeighborsClassifier

ex6_model = KNeighborsClassifier(1).fit(x_train, y_train)
```

Compute training and test accuracy.

```
[ ]: print(f'Training accuracy: {ex6_model.score(x_train, y_train)*100 : .02f}%')
print(f'Test accuracy: {ex6_model.score(x_test, y_test)*100 : .02f}%')
```

Training accuracy: 100.00%

Test accuracy: 92.98%

As we saw in the case of KNN regression, the KNN method is perfect (100% accuracy) on training data with  $K = 1$  by construction. It still achieves high accuracy (but not 100%!) on the test data. The accuracy is comparable to that achieved by the logistic regression.

Repeat for  $K \in \{2, 3, \dots, 9\}$ .

```
[ ]: ex6_models = []
      for k in range(1, 12):
          ex6_models.append(KNeighborsClassifier(k).fit(x_train, y_train))

      for k in range(len(ex6_models)):
          print(f'K = {k + 1}')
          print(f'Training accuracy: {ex6_models[k].score(x_train, y_train)*100 : .
↵02f}%')
          print(f'Test accuracy: {ex6_models[k].score(x_test, y_test)*100 : .02f}%')
          print()
```

K = 1  
 Training accuracy: 100.00%  
 Test accuracy: 92.98%

K = 2  
 Training accuracy: 94.29%  
 Test accuracy: 86.84%

K = 3  
 Training accuracy: 93.85%  
 Test accuracy: 90.35%

K = 4  
 Training accuracy: 92.53%  
 Test accuracy: 91.23%

K = 5  
 Training accuracy: 92.75%  
 Test accuracy: 92.98%

K = 6  
 Training accuracy: 91.87%  
 Test accuracy: 92.98%

K = 7  
 Training accuracy: 92.75%  
 Test accuracy: 94.74%

K = 8  
 Training accuracy: 91.65%  
 Test accuracy: 95.61%

K = 9  
 Training accuracy: 92.97%  
 Test accuracy: 95.61%

K = 10



Training accuracy: 91.87%  
Test accuracy: 94.74%

K = 11  
Training accuracy: 92.97%  
Test accuracy: 95.61%

A similar trend to what occurred with KNN regression is observable here: as  $K$  increases, training accuracy tends to decrease while test accuracy tends to increase. The training accuracy seems to level out at around 92% for  $K \geq 4$ , and the test accuracy seems to top out at around 95% for  $K \geq 8$ .

### 1.2.3 Exercise 7

Create one-layer neural network with sigmoid activation function and binary cross-entropy loss function.

```
[ ]: ex7_model = tf.keras.models.Sequential([
    tf.keras.layers.Input(3),
    tf.keras.layers.Dense(1, activation='sigmoid'),
], name='Exercise7Model')

optim = tf.keras.optimizers.Adam(learning_rate=.1)
ex7_model.compile(loss='bce', optimizer=optim, metrics=['accuracy'])
```

Train for 5000 epochs on the training data.

```
[ ]: x_train_t = tf.convert_to_tensor(x_train)
y_train_t = tf.convert_to_tensor(y_train)

ex7_model.fit(
    x_train_t, y_train_t,
    batch_size=len(x_train),
    epochs=5000,
    verbose=0
)
```

```
[ ]: <keras.src.callbacks.History at 0x2b162857010>
```

Read the network parameters.

```
[ ]: print(f'Weights: {ex7_model.layers[0].weights[0].value()}')
print(f'Bias: {ex7_model.layers[0].bias.value()}')
```

```
Weights: [[27.33742 ]
 [10.767428]
 [15.436617]]
Bias: [-19.742088]
```

These parameters are nearly the same as those obtained using `LogisticRegression`.

```
[ ]: print(f'Training accuracy: {ex7_model.evaluate(x_train_t, y_train_t,
↳ verbose=0)[1]*100 : .02f}%')
print(f'Test accuracy: {ex7_model.evaluate(x_test, y_test, verbose=0)[1]*100 : .
↳ 02f}%')
```

Training accuracy: 92.75%

Test accuracy: 94.74%

The training and test accuracies are, of course, the same as those obtained from LogisticRegression in Exercise 5, as this model is virtually the same model – the parameters were optimized by a different method, but we have already seen that the end result was nearly the same.

Train with 12000 epochs.

```
[ ]: ex7_model_long = tf.keras.models.Sequential([
    tf.keras.layers.Input(3),
    tf.keras.layers.Dense(1, activation='sigmoid'),
], name='Exercise7Model_long')

optim = tf.keras.optimizers.Adam(learning_rate=.1)
ex7_model_long.compile(loss='bce', optimizer=optim, metrics=['accuracy'])

ex7_model_long.fit(
    x_train_t, y_train_t,
    batch_size=len(x_train),
    epochs=12000,
    verbose=0
)
```

```
[ ]: <keras.src.callbacks.History at 0x2b166bb6850>
```

Show learned parameters and training and test accuracies.

```
[ ]: print(f'Weights: {ex7_model_long.layers[0].weights[0].value()}')
print(f'Bias: {ex7_model_long.layers[0].bias.value()}')

print(f'Training accuracy: {ex7_model_long.evaluate(x_train_t, y_train_t,
↳ verbose=0)[1]*100 : .02f}%')
print(f'Test accuracy: {ex7_model_long.evaluate(x_test, y_test,
↳ verbose=0)[1]*100 : .02f}%')
```

Weights: [[27.33744 ]

[10.767424]

[15.436618]]

Bias: [-19.74212]

Training accuracy: 92.75%

Test accuracy: 94.74%

The parameters and corresponding training and test accuracies are the same as those obtained

when training for only 5000 epochs; evidently, the optimization converges before 5000 epochs have passed.

### 1.2.4 Exercise 8

Create model with three hidden layers of size 128, 64, and 32, and one output layer.

```
[ ]: ex8_model = tf.keras.models.Sequential([
    tf.keras.layers.Input(3),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
], 'Exercise8Model')

optim = tf.keras.optimizers.Adam(learning_rate=.1)
ex8_model.compile(loss='bce', optimizer=optim, metrics=['accuracy'])
```

Train for 5000 epochs on the training data.

```
[ ]: ex8_model.fit(
    x_train_t, y_train_t,
    batch_size=len(x_train),
    epochs=5000,
    verbose=0
)
```

```
[ ]: <keras.src.callbacks.History at 0x2b16896c050>
```

Compute training and test accuracies.

```
[ ]: print(f'Training accuracy: {ex8_model.evaluate(x_train_t, y_train_t,
    ↪ verbose=0)[1]*100 : .02f}%')
print(f'Test accuracy: {ex8_model.evaluate(x_test, y_test, verbose=0)[1]*100 : .
    ↪ 02f}%')
```

Training accuracy: 94.07%

Test accuracy: 94.74%

This model achieves slightly higher training and test accuracy than the model from Exercise 7. It has a lot more parameters, though:

```
[ ]: ex7_model.summary()
```

Model: "Exercise7Model"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 1)	4

```

=====
Total params: 4 (16.00 Byte)
Trainable params: 4 (16.00 Byte)
Non-trainable params: 0 (0.00 Byte)
-----

```

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 1)	4

```

=====
Total params: 4 (16.00 Byte)
Trainable params: 4 (16.00 Byte)
Non-trainable params: 0 (0.00 Byte)
-----

```

```
[ ]: ex8_model.summary()
```

```
Model: "Exercise8Model"
```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 128)	512
dense_8 (Dense)	(None, 64)	8256
dense_9 (Dense)	(None, 32)	2080
dense_10 (Dense)	(None, 1)	33

```

=====
Total params: 10881 (42.50 KB)
Trainable params: 10881 (42.50 KB)
Non-trainable params: 0 (0.00 Byte)
-----

```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 128)	512
dense_8 (Dense)	(None, 64)	8256
dense_9 (Dense)	(None, 32)	2080
dense_10 (Dense)	(None, 1)	33

```

=====
Total params: 10881 (42.50 KB)

```

Trainable params: 10881 (42.50 KB)  
Non-trainable params: 0 (0.00 Byte)

---