

Math 6108 Homework 6

Jacob Hauck

October 7, 2024

Question 1.

Question 2.

Let $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n \in \mathbb{R}^n$ be orthonormal. Let $A \in \mathbb{R}^n$. If $A\mathbf{x}_1, A\mathbf{x}_2, \dots, A\mathbf{x}_n$ are also orthonormal, then A is orthogonal.

Proof. Let $X = [\mathbf{x}_1 \ \mathbf{x}_2 \ \cdots \ \mathbf{x}_n]$, and let $B = [A\mathbf{x}_1 \ A\mathbf{x}_2 \ \cdots \ A\mathbf{x}_n]$. Then $B = AX$. Since the columns of B and X are orthonormal, they are both orthogonal matrices. Therefore,

$$A = BX^T.$$

Since $(BX^T)^T(BX^T) = XB^T BX^T = I$ and $(BX^T)(BX^T)^T = BX^T XB^T = I$ by the orthogonality of B and X , it follows that A is invertible with $A^{-1} = (BX^T)^T = A^T$. This implies that A is orthogonal. \square

Question 3.

The algorithm for the Gram-Schmidt process is described in Algorithm 1. An implementation in Python is provided in Listing 1. We note that this implementation detects linear dependence of the columns of A as a part of the Gram-Schmidt process by checking if the produced orthogonal vectors are zero (well, almost zero, to account for numerical rounding error). This is possible because the columns of A are linearly dependent if and only the Gram-Schmidt process produces a zero vector at some point. This is easy to prove.

For $j < i$, each \mathbf{b}_j is a linear combination of the first j columns of A . We can prove this by induction. For the base case, $\mathbf{b}_1 = \|\mathbf{a}_1\|^{-1}\mathbf{a}_1$. For some $1 \leq k < i - 1$, suppose for induction that $\mathbf{b}_j = \sum_{m=1}^j c_{jm}\mathbf{a}_m$ for $1 \leq j \leq k$ and some constants c_{jm} . Then

$$\mathbf{b}_{k+1} = \mathbf{a}_{k+1} - \sum_{p=1}^k \langle \mathbf{b}_p, \mathbf{a}_{k+1} \rangle \sum_{m=1}^p c_{pm}\mathbf{a}_m$$

which completes the proof by induction.

Suppose that $\mathbf{b}_i = \mathbf{0}$. Then

$$\mathbf{0} = \mathbf{b}_i = \mathbf{a}_i - \sum_{p=1}^{i-1} \langle \mathbf{b}_p, \mathbf{a}_i \rangle \sum_{m=1}^p c_{pm}\mathbf{a}_m.$$

The coefficient of \mathbf{a}_i is non-zero, so a non-trivial linear combination of the columns of A is $\mathbf{0}$, meaning that the columns of A are linearly dependent.

Conversely, if the columns of A are linearly dependent, then there exists c_1, \dots, c_m not all equal to zero such that

$$\sum_{i=1}^m c_i \mathbf{a}_i = \mathbf{0}.$$

Let k be the largest integer such that $c_k \neq 0$. Then

$$\mathbf{0} = \sum_{i=1}^k c_i \mathbf{a}_i = c_k \mathbf{b}_k + \sum_{p=1}^{k-1} \langle \mathbf{b}_p, \mathbf{a}_k \rangle \mathbf{b}_p + \sum_{i=1}^{k-1} \left(c_i \mathbf{b}_i + c_i \sum_{p=1}^{i-1} \langle \mathbf{b}_p, \mathbf{a}_i \rangle \mathbf{b}_p \right).$$

The coefficient of \mathbf{b}_k is nonzero, so it follows that the columns of B are linearly dependent. Since the columns of B are also orthogonal because of the Gram-Schmidt process, one of them must be zero.

The command `python -m gs` can be used to run the tests, which verify that the function works across a range of input types that cover every code path. The output from running these tests is given in Listing 2.

Algorithm 1: Gram-Schmidt Orthogonalization

Input: Matrix $A \in \mathbb{R}^{n \times m}$ with linearly independent columns $\mathbf{a}_1, \dots, \mathbf{a}_m \in \mathbb{R}^n$

Output: Matrix $B \in \mathbb{R}^{n \times m}$, whose columns $\mathbf{b}_1, \dots, \mathbf{b}_m \in \mathbb{R}^n$ are the orthogonal vectors obtained by applying the Gram-Schmidt process to the columns of A

```

1  $c \leftarrow 1$ ;
2 repeat
3    $\mathbf{b}_c \leftarrow \mathbf{a}_c - \sum_{p=1}^{c-1} \langle \mathbf{b}_p, \mathbf{a}_c \rangle \mathbf{b}_p$ ;           // Sum is 0 by convention if  $c = 1$ 
4    $\mathbf{b}_c \leftarrow \frac{\mathbf{b}_c}{\|\mathbf{b}_c\|}$ 
5 until  $c = n$ ;
```

Listing 1: Python implementation of the Gram-Schmidt process

```

1 import numpy as np
2
3
4 class LinearDependenceError(BaseException):
5     """Exception class that is raised to indicate linearly dependent input vectors
6     """
7     pass
8
9
10 def gram_schmidt(a, eps_d=1e-10):
11     """
12     Perform the Gram-Schmidt orthogonalization process on the columns of
13     a matrix, returning orthogonalized vectors as the columns of a new matrix.
14
15     :param a: n x m matrix with linearly independent columns. Raises
16               SingularMatrixError if columns of a are linearly dependent or almost
17               linearly dependent (see eps_d).
18     :param eps_d: Tolerance for approximate linear independence (minimum norm of
19                   the computed orthogonal columns). Default = 10^{-10}
20     :return: n x m matrix b whose columns are orthogonal (orthonormal, if
21              normalize == True) vectors obtained by performing Gram-Schmidt
22              orthogonalization on the columns of a.
23     """
24
25     # ==== Input Validation ====
26
27     # Ensure input has the correct data type
```

```

28 a = np.array(a, dtype=float)
29 assert len(a.shape) == 2
30
31 # Early check for linear dependence
32 if a.shape[1] > a.shape[0]:
33     raise LinearDependenceError('Matrix has linearly dependent columns '
34                                 '(more columns than rows)')
35
36 # ==== Run Gram-Schmidt process ====
37
38 # Initialization
39
40 # The first step for each column is copying the corresponding column from a,
41 # so we initialize the output equal to a. Since we already copied the input
42 # with np.array(), we can use that memory for our output matrix
43 b = a
44
45 # Normalize the first column of b
46 norm = np.linalg.norm(b[:, 0])
47
48 # Check for approximate linear dependence before possible divide-by-zero
49 if norm < eps_d:
50     raise LinearDependenceError('Matrix has linearly dependent or almost '
51                                 'linearly dependent columns because first '
52                                 'column is almost 0')
53 b[:, 0] /= norm
54
55 # Iteration
56 for col in range(1, a.shape[1]):
57     # Recall that b[:, col] == a[:, col] because of initialization
58
59     # Subtract out previous orthonormal columns
60     b[:, col] -= b[:, :col] @ (b[:, :col].T @ b[:, col])
61
62     # Normalize new column
63     norm = np.linalg.norm(b[:, col])
64
65     # Check for approximate linear dependence before possible divide-by-zero
66     if norm < eps_d:
67         raise LinearDependenceError('Aborting orthogonalization; matrix has '
68                                     'linearly dependent or almost linearly '
69                                     'dependent columns')
70     b[:, col] /= norm
71
72 # Return orthogonal columns
73 return b
74
75
76 # Test example
77 if __name__ == '__main__':
78     # Set RNG seed for reproducible results

```

```
79     np.random.seed(2024)
80
81     print('Test 1: random square matrix')
82     a = np.random.random((5, 5))
83     print('Input matrix')
84     print(a)
85     print()
86     print('Orthonormalized matrix')
87     b = gram_schmidt(a)
88     print(b)
89     print()
90     print('Implementation worked?', np.allclose(b.T @ b, np.eye(5)))
91     print()
92
93     print('Test 2: random non-square matrix')
94     a = np.random.random((5, 3))
95     print('Input matrix')
96     print(a)
97     print()
98     print('Orthonormalized matrix')
99     b = gram_schmidt(a)
100    print(b)
101    print()
102    print('Implementation worked?', np.allclose(b.T @ b, np.eye(3)))
103    print()
104
105    print('Test 3: random matrix with too many columns')
106    a = np.random.random((3, 5))
107    print('Input matrix')
108    print(a)
109    print()
110    try:
111        gram_schmidt(a) # should raise an error
112    except LinearDependenceError as e:
113        print(e)
114    print()
115
116    print('Test 4: matrix with first column 0')
117    a = np.array([
118        [0, 1, 2],
119        [0, 3, 4],
120        [0, 5, 6]
121    ])
122    print('Input matrix')
123    print(a)
124    print()
125    try:
126        gram_schmidt(a) # should raise an error
127    except LinearDependenceError as e:
128        print(e)
129    print()
```

```

130
131     print('Test 5: singular matrix')
132     a = np.array([
133         [1, 2, -1],
134         [2, 5, -3],
135         [3, 3, 0]
136     ])
137     print('Input matrix')
138     print(a)
139     print()
140     try:
141         gram_schmidt(a) # should raise an error
142     except LinearDependenceError as e:
143         print(e)

```

Listing 2: Output for test cases

```

1 >python -m gs
2 Test 1: random square matrix
3 Input matrix
4 [[0.58801452 0.69910875 0.18815196 0.04380856 0.20501895]
5  [0.10606287 0.72724014 0.67940052 0.4738457 0.44829582]
6  [0.01910695 0.75259834 0.60244854 0.96177758 0.66436865]
7  [0.60662962 0.44915131 0.22535416 0.6701743 0.73576659]
8  [0.25799564 0.09554215 0.96090974 0.25176729 0.28216512]]
9
10 Orthonormalized matrix
11 [[ 0.66075857  0.10599106 -0.32621454 -0.56072372 -0.36240445]
12  [ 0.11918405  0.62410254  0.17611003 -0.29186203  0.69288743]
13  [ 0.02147069  0.73799494  0.08061366  0.44264935 -0.50245942]
14  [ 0.68167657 -0.1644837  -0.10793795  0.62668109  0.32230789]
15  [ 0.28991262 -0.16604366  0.91892338 -0.10834125 -0.17950537]]
16
17 Implementation worked? True
18
19 Test 2: random non-square matrix
20 Input matrix
21 [[0.76825393 0.7979234 0.5440372 ]
22  [0.38270763 0.38165095 0.28582739]
23  [0.74026815 0.23898683 0.4377217 ]
24  [0.8835387 0.28928114 0.78450686]
25  [0.75895366 0.41778538 0.22576877]]
26
27 Orthonormalized matrix
28 [[ 0.47270878  0.73926285  0.17771326]
29  [ 0.23548107  0.33566518  0.12907525]
30  [ 0.45548905 -0.37843189 -0.14903895]
31  [ 0.54364382 -0.44335429  0.57756706]
32  [ 0.46698629 -0.03233593 -0.77198527]]
33
34 Implementation worked? True
35

```

```
36 Test 3: random matrix with too many columns
37 Input matrix
38 [[0.42009814 0.06436369 0.59643269 0.83732372 0.89248639]
39 [0.20052744 0.50239523 0.89538184 0.25592093 0.86723234]
40 [0.01648793 0.55249695 0.52790539 0.92335039 0.24594844]]
41
42 Matrix has linearly dependent columns (more columns than rows)
43
44 Test 4: matrix with first column 0
45 Input matrix
46 [[0 1 2]
47 [0 3 4]
48 [0 5 6]]
49
50 Matrix has linearly dependent or almost linearly dependent columns because first
    ↪ column is almost 0
51
52 Test 5: singular matrix
53 Input matrix
54 [[ 1  2 -1]
55 [ 2  5 -3]
56 [ 3  3  0]]
57
58 Aborting orthogonalization; matrix has linearly dependent or almost linearly dependent
    ↪ columns
```